

Code Experiment

April 16, 2020

1 Base Model

1.1 Importing Necessary Dependencies

```
[1]: # Python Standard Libraries for importing data from binary file
import os.path #for accessing the file path
import struct  #for unpacking the binary data

import time    #for calculating time
import math

#core packages
import numpy as np
import matplotlib.pyplot as plt

#custom module
from dataPrep import retrieve_data, sample_origDataset , dev_test_split, \
    ↪prep_dataset, visualize_orig

np.random.seed(1)
%matplotlib inline
```

1.2 Dataset Preparation

Retriving Dataset

```
[2]: #retriving the data
train_x_orig, train_y_orig = retrieve_data(dataset="training-set")
test_x_temp, test_y_temp = retrieve_data(dataset="test-set")

#displaying the retrival info
print("Data\t\t", "Datatype\t", "Shape")
print("=====")
print("Training Set Images:\t" + str(type(train_x_orig))+"\t", str(train_x_orig.
    ↪shape))
```

```

print("Training Set Labels:\t" + str(type(train_y_orig))+ "\t",str(train_y_orig.
↪shape))
print("Test Set Images:\t" + str(type(test_x_temp))+ "\t",str(test_x_temp.shape))
print("Test Set Labels:\t" + str(type(test_y_temp))+ "\t",str(test_y_temp.shape))
print("=====")

```

Data	Datatype	Shape
Training Set Images:	<class 'numpy.ndarray'>	(60000, 28, 28)
Training Set Labels:	<class 'numpy.ndarray'>	(60000, 1)
Test Set Images:	<class 'numpy.ndarray'>	(10000, 28, 28)
Test Set Labels:	<class 'numpy.ndarray'>	(10000, 1)

Sampling the Dataset for Model Experiment

```

[3]: train_Vol,train_x_sample, train_y_sample =
↪sample_origDataset(train_x_orig,train_y_orig,dataVol= 100)
test_Vol,test_x_sample,test_y_sample =
↪sample_origDataset(test_x_temp,test_y_temp,dataVol= 100)

print("Data\t\t\t","Complete Data Shape\t","Sample Data Shape\t","Sample Size")
print("=====")
print("Training Set Images:\t"+ str(train_x_orig.shape)+ "\t\t"+
↪str(train_x_sample.shape)+ "\t\t"+str(train_Vol)+"%")
print("Training Set Labels:\t"+ str(train_y_orig.shape)+ "\t\t"+
↪str(train_y_sample.shape))
print("Test Set Images:\t"+str(test_x_temp.shape)+ "\t\t"+ str(test_x_sample.
↪shape)+ "\t\t"+str(test_Vol)+"%")
print("Test Set Labels:\t"+str(test_y_temp.shape)+ "\t\t"+ str(test_y_sample.
↪shape))
print("=====")

```

Data Size	Complete Data Shape	Sample Data Shape	Sample Size
Training Set Images:	(60000, 28, 28)	(60000, 28, 28)	100%
Training Set Labels:	(60000, 1)	(60000, 1)	
Test Set Images:	(10000, 28, 28)	(10000, 28, 28)	100%
Test Set Labels:	(10000, 1)	(10000, 1)	

Splitting Dev-Test Set

```
[4]: dev_x_orig,dev_y_orig,test_x_orig,test_y_orig =
      ↪dev_test_split(test_x_sample,test_y_sample)

print("Data\t\t\t","Shape")
print("=====")
print("Dev Set Images:\t\t",str(dev_x_orig.shape))
print("Dev Set Labels:\t\t",str(dev_y_orig.shape))
print("Test Set Images:\t",str(test_x_orig.shape))
print("Test Set Labels:\t",str(test_y_orig.shape))
print("=====")
```

Data	Shape
=====	
Dev Set Images:	(5000, 28, 28)
Dev Set Labels:	(5000, 1)
Test Set Images:	(5000, 28, 28)
Test Set Labels:	(5000, 1)
=====	

Preparing the Dataset (Flattening and Normalizing)

```
[5]: train_x_norm,train_y_encoded, dev_x_norm,dev_y_encoded, test_x_norm,
      ↪test_y_encoded = prep_dataset(train_x_sample, train_y_sample, dev_x_orig,
      ↪dev_y_orig, test_x_orig, test_y_orig)

print("Data\t\t\t","Before Processing\t","After Processing")
print("=====")
print("Training Set Images:\t" + str(train_x_orig.shape)+"\t\t"+
      ↪str(train_x_norm.shape))
print("Training Set Labels:\t" + str(train_y_orig.shape)+"\t\t"+
      ↪str(train_y_encoded.shape))
print("Dev Set Images:\t\t" + str(dev_x_orig.shape)+"\t\t"+ str(dev_x_norm.
      ↪shape))
print("Dev Set Labels:\t\t" + str(dev_y_orig.shape)+"\t\t"+ str(dev_y_encoded.
      ↪shape))
print("Test Set Images:\t" + str(test_x_orig.shape)+"\t\t"+ str(test_x_norm.
      ↪shape))
print("Test Set Labels:\t" + str(test_y_orig.shape)+"\t\t"+ str(test_y_encoded.
      ↪shape))
print("=====")
```

Data	Before Processing	After Processing
=====		
Training Set Images:	(60000, 28, 28)	(784, 60000)
Training Set Labels:	(60000, 1)	(11, 60000)
Dev Set Images:	(5000, 28, 28)	(784, 5000)
Dev Set Labels:	(5000, 1)	(11, 5000)
Test Set Images:	(5000, 28, 28)	(784, 5000)

Test Set Labels: (5000, 1) (11, 5000)
=====

Creating Minibatches

```
[6]: def rand_mini_batches(X, Y, mini_batch_size = 64, seed=1):

    np.random.seed(seed)
    m = X.shape[1]                # number of training examples
    mini_batches = []

    # Shuffle (X, Y)
    permutation = list(np.random.permutation(m))
    shuffled_X = X[:, permutation]
    shuffled_Y = Y[:, permutation].reshape((11,m))

    # Partition (shuffled_X, shuffled_Y) except for the last batch
    num_complete_minibatches = math.floor(m/mini_batch_size) # number of mini_
    ↳ batches of size mini_batch_size
    for k in range(0, num_complete_minibatches):
        mini_batch_X = shuffled_X[:, k * mini_batch_size :
    ↳ (k+1)*mini_batch_size]
        mini_batch_Y = shuffled_Y[:, k * mini_batch_size :
    ↳ (k+1)*mini_batch_size]
        mini_batch = (mini_batch_X, mini_batch_Y)
        mini_batches.append(mini_batch)

    # Last batch (last mini-batch < mini_batch_size)
    if m % mini_batch_size != 0:
        mini_batch_X = shuffled_X[:, num_complete_minibatches * mini_batch_size
    ↳ : m]
        mini_batch_Y = shuffled_Y[:, num_complete_minibatches * mini_batch_size
    ↳ : m]
        mini_batch = (mini_batch_X, mini_batch_Y)
        mini_batches.append(mini_batch)

    return mini_batches
```

```
[7]: mini_batches = rand_mini_batches(train_x_norm, train_y_encoded, mini_batch_size
    ↳ = 128)
minibatch_length = len(mini_batches)
print("Total Minibatches:\t"+str(minibatch_length))
print("\nMinibatches\t\tShape")
print("=====")
print ("1st mini_batch_X:\t" + str(mini_batches[0][0].shape))
print ("2nd mini_batch_X:\t" + str(mini_batches[1][0].shape))
```

```

print (str(minibatch_length - 1) + "th mini_batch_X:\t" +
      ↪str(mini_batches[-2][0].shape))
print (str(minibatch_length) + "th mini_batch_X:\t" + str(mini_batches[-1][0].
      ↪shape))

print ("\n1st mini_batch_Y:\t" + str(mini_batches[0][1].shape))
print ("2nd mini_batch_Y:\t" + str(mini_batches[1][1].shape))
print (str(minibatch_length - 1) + "th mini_batch_Y:\t" +
      ↪str(mini_batches[-2][1].shape))
print (str(minibatch_length) + "th mini_batch_Y:\t" + str(mini_batches[-1][1].
      ↪shape))

```

Total Minibatches: 469

Minibatches	Shape
=====	
1st mini_batch_X:	(784, 128)
2nd mini_batch_X:	(784, 128)
468th mini_batch_X:	(784, 128)
469th mini_batch_X:	(784, 96)
1st mini_batch_Y:	(11, 128)
2nd mini_batch_Y:	(11, 128)
468th mini_batch_Y:	(11, 128)
469th mini_batch_Y:	(11, 96)

1.3 Utility Functions

ReLU Function and Its derivative

```

[8]: def relu(Z):
      A = np.maximum(0.0,Z)

      cache = Z
      assert(A.shape == Z.shape)
      return A, cache

[9]: def relu_grad(dA, cache):
      Z = cache
      dZ = np.array(dA, copy=True) # just converting dz to a correct object.

      dZ[Z < 0] = 0

      assert(dZ.shape == Z.shape)
      return dZ

```

Softmax Function and its derivative

```
[10]: def softmax(Z):  
    shift = Z - np.max(Z) #Avoiding underflow or overflow errors due to  
    →floating point instability in softmax  
    t = np.exp(shift)  
    A = np.divide(t,np.sum(t,axis = 0))  
  
    cache = Z  
    assert(A.shape == Z.shape)  
    return A, cache
```

Learning rate decay

```
[12]: def decay_learning_rate(alpha_prev, epoch, decay_rate = 1 ):  
  
    alpha = (1/(1 + decay_rate * epoch)) * alpha_prev  
  
    return alpha
```

1.4 Deep Learning Model

1.4.1 1. Creating NN Architecture

initializing layers

```
[13]: def init_layers():  
    layers_dim = [784,800,300,11]  
    return layers_dim
```

Initializing Parameters

- Random initialization

```
[14]: def init_params_random(layers_dim):  
  
    L = len(layers_dim)  
    params = {}  
  
    for l in range(1,L):  
        params['W' + str(l)] = np.random.randn(layers_dim[l],layers_dim[l-1])  
        →*0.01  
        params['b' + str(l)] = np.zeros((layers_dim[l],1))  
  
        assert(params['W' + str(l)].shape == (layers_dim[l],layers_dim[l-1]))  
        assert(params['b' + str(l)].shape == (layers_dim[l],1))  
    return params
```

- He-initialization

```
[15]: def init_params_he(layers_dim):

    L = len(layers_dim)
    params = {}

    for l in range(1,L):
        params['W' + str(l)] = np.random.randn(layers_dim[l],layers_dim[l-1]) *
        ↪ np.sqrt(np.divide(2,layers_dim[l-1])) # He - initialization
        params['b' + str(l)] = np.zeros((layers_dim[l],1))

        assert(params['W' + str(l)].shape == (layers_dim[l],layers_dim[l-1]))
        assert(params['b' + str(l)].shape == (layers_dim[l],1))
    return params
```

Initializing Hyper Parameters

```
[16]: def init_hyperParams(alpha = 0.001, num_epoch = 2000, mini_batch_size = 128,
        ↪ beta1 = 0.9, beta2 = 0.999, epsilon = 1e-8):
    hyperParams = {}
    hyperParams['learning_rate'] = alpha
    hyperParams['num_epoch'] = num_epoch
    hyperParams['mini_batch_size'] = mini_batch_size
    hyperParams['beta1'] = beta1
    hyperParams['beta2'] = beta2
    hyperParams['epsilon'] = epsilon

    return hyperParams
```

1.4.2 2. Forward Propagation

Calculating sum of product of inputs and weights (Z) for individual layer

```
[17]: def forward_sum(A,W,b):

    Z = np.dot(W,A) + b

    cache = (A,W,b)
    assert(Z.shape == (W.shape[0],Z.shape[1]))

    return Z, cache
```

Calculating Activation for individual Layer

```
[18]: def forward_activation(A,W,b,activation):

    if activation == 'relu':
        Z, sum_cache = forward_sum(A,W,b)
        A, activation_cache = relu(Z)

    if activation == 'softmax':
        Z, sum_cache = forward_sum(A,W,b)
        A, activation_cache = softmax(Z)

    cache = (sum_cache,activation_cache)
    assert(A.shape == Z.shape)

    return A, cache
```

Complete Forward Propagation for L layers

```
[19]: def forward_prop(X,parameters):
    caches = []
    A = X
    L = len(parameters) // 2
    for l in range(1, L):
        A_prev = A
        A, cache = forward_activation(A_prev,parameters['W' +
→str(l)],parameters['b' + str(l)],activation='relu')
        caches.append(cache)

    AL, cache = forward_activation(A,parameters['W' + str(L)],parameters['b' +
→str(L)],activation='softmax')
    caches.append(cache)

    assert(AL.shape == (11,X.shape[1]))

    return AL,caches
```

1.4.3 3. Cost Function

```
[20]: def compute_cost(AL,Y):
    m = Y.shape[1]

    cost = -(1./m) * np.sum(np.sum(np.multiply(Y,np.log(AL)), axis =
→0,keepdims=True))
```



```

    cost = np.squeeze(cost)      # Making sure your cost's shape is not
    ↪ returned as ndarray
    assert(cost.shape == ())

    return cost

```

1.4.4 4. Backward Propagation

Calculating Gradients for individual Layer

```

[21]: def backward_grad(dZ, cache):
    A_prev, W, b = cache
    m = A_prev.shape[1]

    dW = (1/m) * np.dot(dZ,A_prev.T)
    db = (1/m) * np.sum(dZ, axis = 1, keepdims=True )
    dA_prev = np.dot(W.T, dZ)

    assert (dA_prev.shape == A_prev.shape)
    assert (dW.shape == W.shape)
    assert (db.shape == b.shape)

    return dA_prev, dW, db

```

Calculating Backward Activation for individual layer

```

[22]: def backward_activation(dA,cache,activation):
    sum_cache, activation_cache = cache

    if activation == "relu":
        dZ = relu_grad(dA,activation_cache)
        dA_prev, dW, db = backward_grad(dZ, sum_cache)

    elif activation == "softmax":
        dZ = dA
        dA_prev, dW, db = backward_grad(dA, sum_cache)

    return dA_prev, dW, db

```

Complete Backward Propagation for L layers

```
[23]: def backward_prop(AL, Y, caches):
    grads = {}
    L = len(caches) # the number of layers
    m = AL.shape[1]
    Y = Y.reshape(AL.shape) # after this line, Y is the same shape as AL

    dA = np.subtract(AL, Y)

    current_cache = caches[L-1]
    grads["dA" + str(L-1)], grads["dW" + str(L)], grads["db" + str(L)] = backward_activation(dA, current_cache, activation = 'softmax')

    for l in reversed(range(L-1)):
        current_cache = caches[l]
        dA_prev_temp, dW_temp, db_temp = backward_activation(grads["dA" + str(l+1)], current_cache, activation = 'relu')
        grads["dA" + str(l)] = dA_prev_temp
        grads["dW" + str(l+1)] = dW_temp
        grads["db" + str(l+1)] = db_temp

    return grads
```

1.4.5 5. Update Parameters

- normal update of parameters

```
[24]: def update_parameters(parameters, grads, learning_rate):
    L = len(parameters) // 2
    for l in range(L):
        parameters["W" + str(l+1)] = parameters["W" + str(l+1)] - (learning_rate * grads["dW" + str(l+1)])
        parameters["b" + str(l+1)] = parameters["b" + str(l+1)] - (learning_rate * grads["db" + str(l+1)])
    return parameters
```

- update parameters with adam

```
[25]: #initialize adam

def initialize_adam(parameters) :

    L = len(parameters) // 2
    v = {}
    s = {}

    for l in range(L):
```

```

v["dW" + str(l+1)] = np.zeros(parameters["W" + str(l+1)].shape)
v["db" + str(l+1)] = np.zeros(parameters["b" + str(l+1)].shape)
s["dW" + str(l+1)] = np.zeros(parameters["W" + str(l+1)].shape)
s["db" + str(l+1)] = np.zeros(parameters["b" + str(l+1)].shape)

return v, s

```

```

[26]: #update with adam
def update_parameters_adam(parameters, grads, learning_rate, v, s, t, beta1 = 0.9, beta2 = 0.999, epsilon = 1e-8):

    L = len(parameters) // 2
    v_corrected = {}
    s_corrected = {}

    for l in range(L):
        # Moving average of the gradients.
        v["dW" + str(l+1)] = np.add(beta1 * v["dW" + str(l+1)], (1 - beta1) *
        ↪ grads["dW" + str(l+1)])
        v["db" + str(l+1)] = np.add(beta1 * v["db" + str(l+1)], (1 - beta1) *
        ↪ grads["db" + str(l+1)])

        # Compute bias-corrected first moment estimate.
        v_corrected["dW" + str(l+1)] = np.divide(v["dW" + str(l+1)], (1 - np.
        ↪ power(beta1, t)))
        v_corrected["db" + str(l+1)] = np.divide(v["db" + str(l+1)], (1 - np.
        ↪ power(beta1, t)))

        # Moving average of the squared gradients.
        s["dW" + str(l+1)] = np.add(beta2 * s["dW" + str(l+1)], (1 - beta2) *
        ↪ np.square(grads["dW" + str(l+1)]))
        s["db" + str(l+1)] = np.add(beta2 * s["db" + str(l+1)], (1 - beta2) *
        ↪ np.square(grads["db" + str(l+1)]))

        # Compute bias-corrected second raw moment estimate.
        s_corrected["dW" + str(l+1)] = np.divide(s["dW" + str(l+1)], (1 - np.
        ↪ power(beta2, t)))
        s_corrected["db" + str(l+1)] = np.divide(s["db" + str(l+1)], (1 - np.
        ↪ power(beta2, t)))

        # Update parameters.
        parameters["W" + str(l+1)] = np.subtract(parameters["W" + str(l+1)],
        ↪ learning_rate * np.divide(v_corrected["dW" + str(l+1)], np.
        ↪ sqrt(s_corrected["dW" + str(l+1)]) + epsilon))

```

```

        parameters["b" + str(l+1)] = np.subtract(parameters["b" + str(l+1)],
↪learning_rate * np.divide(v_corrected["db" + str(l+1)], np.
↪sqrt(s_corrected["db" + str(l+1)]) + epsilon))

    return parameters, v, s

```

1.4.6 6. Prediction

```

[27]: def predict(X,y,parameters):
    m = y.shape[1]
    n = len(parameters) // 2 # number of layers in the neural network

    probas, caches = forward_prop(X, parameters)

    assert(probas.shape == y.shape)

    predicted_labels = np.argmax(probas,axis=0).reshape(1,probas.shape[1])
    predicted_prob = np.max(probas,axis = 0).reshape(1,m)

    Y = np.argmax(y,axis=0).reshape(1,y.shape[1])

    #print results
    true_prediction = np.equal(predicted_labels,Y)

    num_correct_labels = np.sum(true_prediction)
    accuracy = (num_correct_labels/m)

    return predicted_labels, predicted_prob, accuracy

```

Visualizing the costs and accuracy for model analysis

```

[57]: def visualize_results(attr, attr_type):

    plt.plot(np.squeeze(attr))
    if attr_type == 'costs':
        plt.ylabel("cost")
        plt.title("Cost")

    elif attr_type == 'train_accs':
        plt.ylabel("accuracy")
        plt.title("Training Accuracy")
    #     plt.plot(np.squeeze(1 - attr), label = 'loss')

```

```

elif attr_type == 'val_accs':
    plt.ylabel("accuracy")
    plt.title("Validation Accuracy")
#     plt.plot(np.squeeze(1 - attr), label = 'loss')

else:
    raise ValueError("Dataset set must be training or dev or test set")

plt.xlabel('Epochs (per hundreds)')
plt.show()

```

1.4.7 7. Train Model

```

[58]: def train(X_train, Y_train, X_dev, Y_dev, layers_dim, hyperParams, optimizer = '
    ↪adam'):
    #hyper parameters
    learning_rate = hyperParams['learning_rate']
    num_epoch = hyperParams['num_epoch']
    mini_batch_size = hyperParams['mini_batch_size']
    beta1 = hyperParams['beta1']
    beta2 = hyperParams['beta2']
    epsilon = hyperParams['epsilon']

    seed = 1
    m = Y_train.shape[1]
    costs = [] # keep track of epoch cost
    train_accs = [] # keep track of training accuracy
    val_accs = [] # keep track of Validation accuracy

    parameters = init_params_he(layers_dim)

    if optimizer == 'mgd':
        pass
    elif optimizer == 'adam':
        t = 0
        v,s = initialize_adam(parameters)

    #minibatch GD
    for i in range(0, num_epoch):
        seed += 1
        minibatches = rand_mini_batches(X_train, Y_train, mini_batch_size, seed)
        epoch_cost = 0

    # learning rate decay
    # if i % 5 == 0:

```

```

#         learning_rate = decay_learning_rate(learning_rate, i , decay_rate
→ = 0.1)

    for minibatch in minibatches:

        (minibatch_X, minibatch_Y) = minibatch

        AL, caches = forward_prop(minibatch_X, parameters)

        epoch_cost += compute_cost(AL, minibatch_Y) #accumulating the batch
→ costs

        grads = backward_prop(AL, minibatch_Y, caches)

        if optimizer == 'mgd':
            parameters = update_parameters(parameters, grads, learning_rate)
        elif optimizer == 'adam':
            t = t+1
            parameters, v, s = update_parameters_adam(parameters,
→ grads, learning_rate, v, s, t, beta1, beta2, epsilon)

        epoch_cost_avg = epoch_cost / m

        #computing and accumulating training and validation accuracy
        _, train_acc = predict(X_train, Y_train, parameters)
        _, val_acc = predict(X_dev, Y_dev, parameters)
        train_accs.append(train_acc)
        val_accs.append(val_acc)

#         if i % 50 == 0:
            print("\nEpoch: %d == Learning rate: %f"%(i, learning_rate))
            print ("t== Cost: %f || Training acc: %.6f || Val acc: %.6f || Val
→ loss: %.6f"%(epoch_cost_avg, train_acc, val_acc, 1-val_acc))
#         if i % 100 == 0:
            costs.append(epoch_cost_avg)

    visualize_results(costs, attr_type='costs')

    visualize_results(train_accs, attr_type='train_accs')
    visualize_results(val_accs, attr_type='val_accs')

    return parameters

```

1.4.8 Running Model

```
[59]: hyperParams = init_hyperParams(alpha = 0.001, num_epoch = 50, mini_batch_size = 512)
      layers_dim = init_layers()
      parameters = train(train_x_norm, train_y_encoded, dev_x_norm, dev_y_encoded, layers_dim, hyperParams, optimizer = 'adam')
```

```
Epoch: 0 == Learning rate: 0.001000
          == Cost: 0.000613 || Training acc: 0.962467 || Val acc: 0.959200 || Val
loss: 0.040800
```

```
Epoch: 1 == Learning rate: 0.001000
          == Cost: 0.000210 || Training acc: 0.982550 || Val acc: 0.973800 || Val
loss: 0.026200
```

```
Epoch: 2 == Learning rate: 0.001000
          == Cost: 0.000132 || Training acc: 0.987433 || Val acc: 0.976400 || Val
loss: 0.023600
```

```
Epoch: 3 == Learning rate: 0.001000
          == Cost: 0.000086 || Training acc: 0.992950 || Val acc: 0.977600 || Val
loss: 0.022400
```

```
Epoch: 4 == Learning rate: 0.001000
          == Cost: 0.000062 || Training acc: 0.993750 || Val acc: 0.976000 || Val
loss: 0.024000
```

```
Epoch: 5 == Learning rate: 0.001000
          == Cost: 0.000043 || Training acc: 0.997733 || Val acc: 0.980600 || Val
loss: 0.019400
```

```
Epoch: 6 == Learning rate: 0.001000
          == Cost: 0.000030 || Training acc: 0.998333 || Val acc: 0.980800 || Val
loss: 0.019200
```

```
Epoch: 7 == Learning rate: 0.001000
          == Cost: 0.000018 || Training acc: 0.998567 || Val acc: 0.981800 || Val
loss: 0.018200
```

```
Epoch: 8 == Learning rate: 0.001000
          == Cost: 0.000011 || Training acc: 0.999750 || Val acc: 0.982200 || Val
loss: 0.017800
```

```
Epoch: 9 == Learning rate: 0.001000
          == Cost: 0.000008 || Training acc: 0.999367 || Val acc: 0.981400 || Val
```

loss: 0.018600

Epoch: 10 == Learning rate: 0.001000
== Cost: 0.000012 || Training acc: 0.999467 || Val acc: 0.980600 || Val
loss: 0.019400

Epoch: 11 == Learning rate: 0.001000
== Cost: 0.000006 || Training acc: 0.998450 || Val acc: 0.980200 || Val
loss: 0.019800

Epoch: 12 == Learning rate: 0.001000
== Cost: 0.000010 || Training acc: 0.999517 || Val acc: 0.981400 || Val
loss: 0.018600

Epoch: 13 == Learning rate: 0.001000
== Cost: 0.000006 || Training acc: 0.997933 || Val acc: 0.978600 || Val
loss: 0.021400

Epoch: 14 == Learning rate: 0.001000
== Cost: 0.000019 || Training acc: 0.997633 || Val acc: 0.979400 || Val
loss: 0.020600

Epoch: 15 == Learning rate: 0.001000
== Cost: 0.000026 || Training acc: 0.996733 || Val acc: 0.977400 || Val
loss: 0.022600

Epoch: 16 == Learning rate: 0.001000
== Cost: 0.000024 || Training acc: 0.998783 || Val acc: 0.979800 || Val
loss: 0.020200

Epoch: 17 == Learning rate: 0.001000
== Cost: 0.000016 || Training acc: 0.998250 || Val acc: 0.979400 || Val
loss: 0.020600

Epoch: 18 == Learning rate: 0.001000
== Cost: 0.000007 || Training acc: 0.999267 || Val acc: 0.983400 || Val
loss: 0.016600

Epoch: 19 == Learning rate: 0.001000
== Cost: 0.000004 || Training acc: 0.999967 || Val acc: 0.983200 || Val
loss: 0.016800

Epoch: 20 == Learning rate: 0.001000
== Cost: 0.000001 || Training acc: 1.000000 || Val acc: 0.983600 || Val
loss: 0.016400

Epoch: 21 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984000 || Val

loss: 0.016000

Epoch: 22 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.983800 || Val
loss: 0.016200

Epoch: 23 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984000 || Val
loss: 0.016000

Epoch: 24 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984000 || Val
loss: 0.016000

Epoch: 25 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984000 || Val
loss: 0.016000

Epoch: 26 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984000 || Val
loss: 0.016000

Epoch: 27 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val
loss: 0.015600

Epoch: 28 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984200 || Val
loss: 0.015800

Epoch: 29 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val
loss: 0.015600

Epoch: 30 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984600 || Val
loss: 0.015400

Epoch: 31 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val
loss: 0.015600

Epoch: 32 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val
loss: 0.015600

Epoch: 33 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val

loss: 0.015600

Epoch: 34 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984600 || Val
loss: 0.015400

Epoch: 35 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val
loss: 0.015600

Epoch: 36 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 37 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val
loss: 0.015600

Epoch: 38 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984600 || Val
loss: 0.015400

Epoch: 39 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 40 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984600 || Val
loss: 0.015400

Epoch: 41 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 42 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984600 || Val
loss: 0.015400

Epoch: 43 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 44 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 45 == Learning rate: 0.001000
== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984400 || Val

loss: 0.015600

Epoch: 46 == Learning rate: 0.001000

== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 47 == Learning rate: 0.001000


== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200



Epoch: 48 == Learning rate: 0.001000

== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.984800 || Val
loss: 0.015200

Epoch: 49 == Learning rate: 0.001000

== Cost: 0.000000 || Training acc: 1.000000 || Val acc: 0.985000 || Val
loss: 0.015000





```
[60]: predicted_labels_train, prediction_prob_train, train_acc = predict(train_x_norm,
    ↪ train_y_encoded, parameters)
print("\nAccuracy: " + str(train_acc))
print("\nError: \t" + str(1-train_acc))
```

Accuracy: 1.0

Error: 0.0

```
[61]: predicted_labels_dev, prediction_prob_dev, dev_acc = predict(dev_x_norm,
    ↪ dev_y_encoded, parameters)
print("\nAccuracy: " + str(dev_acc))
print("\nError: \t" + str(1-dev_acc))
```

Accuracy: 0.985

Error: 0.0150000000000000013

```
[62]: predicted_labels_test, prediction_prob_test, test_acc = predict(test_x_norm,
    ↪ test_y_encoded, parameters)
print("\nAccuracy: " + str(test_acc))
print("\nError: \t" + str(1-test_acc))
```

Accuracy: 0.9878

Error: 0.012199999999999989

Visualizing Prediction

```
[63]: def visualize_prediction(x_orig, y_orig, predicted_labels, prediction_prob,
    ↪ dataset):
    if(dataset == "training"):
        visual_title = "Sample Training Data Set"
        rng = range(30,40)
    elif(dataset == "dev"):
        visual_title = "Sample Dev Data Set"
        rng = range(110,120)
    elif(dataset == "test"):
        visual_title = "Sample Test Data Set"
        rng = range(110,120)
    else:
        raise ValueError("Dataset set must be training or dev or test set")
    fig, axes = plt.subplots(nrows=2, ncols=5, figsize=(16,8))
    fig.subplots_adjust(hspace=1)
```

```

fig.suptitle(visual_title)

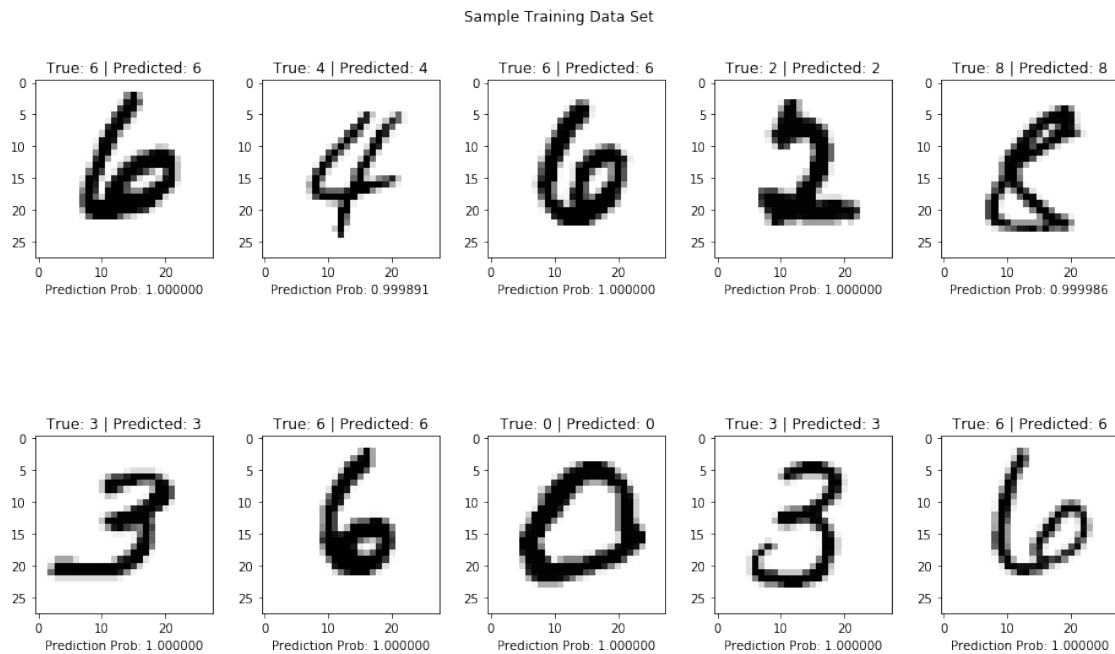
for ax,i in zip(axes.flatten(),rng):
    ax.imshow(x_orig[i].squeeze(),interpolation='nearest', cmap='Greys')
    ax.set(title = "True: " + str(y_orig[0,i])+" | Predicted:␣
↪"+str(predicted_labels[0,i]))
    ax.set(xlabel= "Prediction Prob: %f"%(prediction_prob[0,i]))

```

```

[64]: visualize_prediction(train_x_sample, train_y_sample.T,predicted_labels_train,␣
↪prediction_prob_train,dataset = "training")

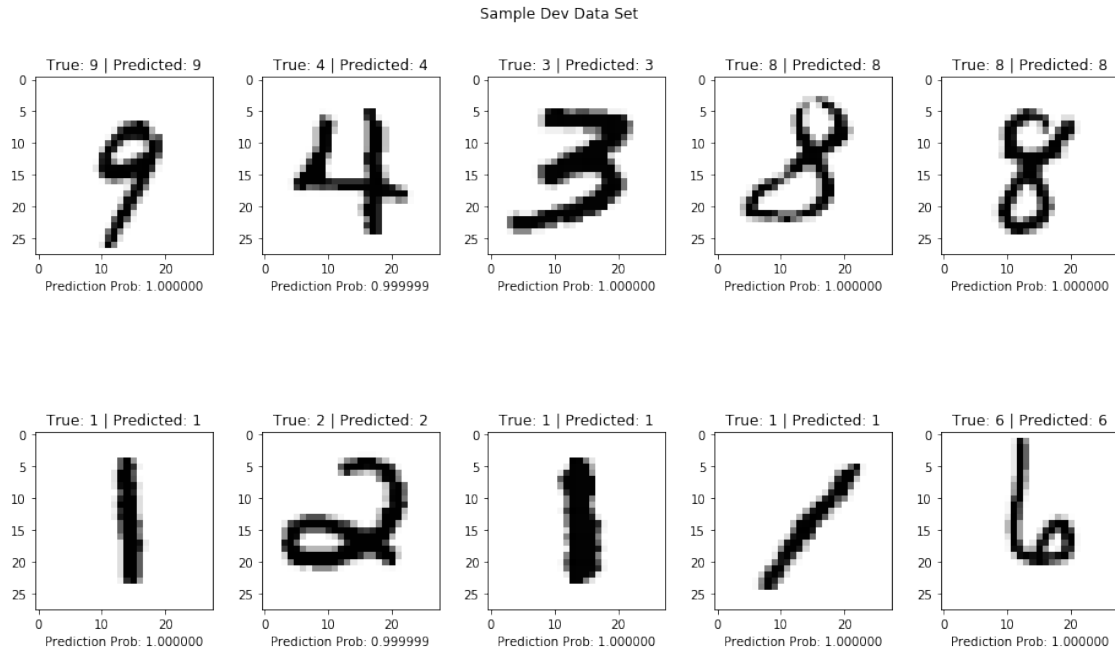
```



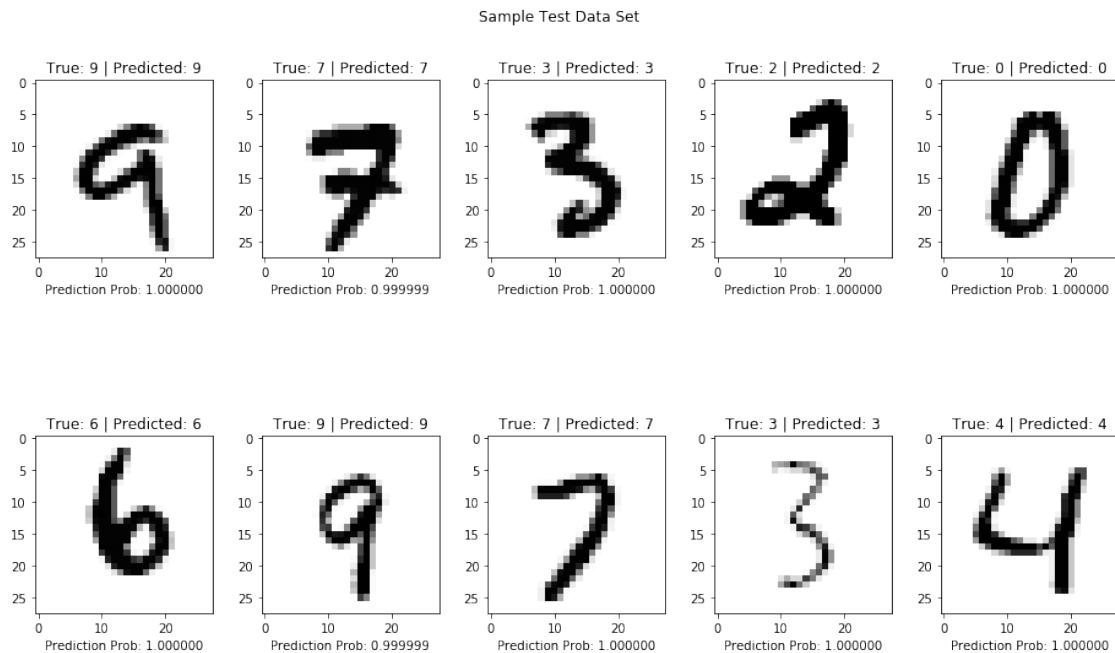
```

[65]: visualize_prediction(dev_x_orig, dev_y_orig.T, predicted_labels_dev,␣
↪prediction_prob_prob,dataset = "dev")

```



```
[66]: visualize_prediction(test_x_orig, test_y_orig.T, predicted_labels_test,
    ↪ prediction_prob_prob, dataset = "test")
```



Visualizing Mislabelled Images in all datasets

```
[67]: def
    ↪ visualize_mislabelled_images(x_orig,y_orig,predicted_labels,prediction_prob,dataset):
    ↪
        true_prediction = np.equal(predicted_labels,y_orig)
        mislabelled_indices = np.asarray(np.where(true_prediction == False))
        print("Total Mislabelled Images: "+str(len(mislabelled_indices[0])))

        if(dataset == "training"):
            visual_title = "Sample Mislabelled Training Images"
        elif(dataset == "dev"):
            visual_title = "Sample Mislabelled Dev Images"
        elif(dataset == "test"):
            visual_title = "Sample Mislabelled Test Images"
        else:
            raise ValueError("Dataset set must be training or dev or test set")


        fig, axes = plt.subplots(nrows=2, ncols=5,figsize=(16,8))
        fig.subplots_adjust(hspace=1)
        fig.suptitle(visual_title)

        for ax,i in zip(axes.flatten(),mislabelled_indices[1]):
            ax.imshow(x_orig[i].squeeze(),interpolation='nearest')
            ax.set(title = "True: " + str(y_orig[0,i])+" | Predicted:
    ↪ "+str(predicted_labels[0,i]))
            ax.set(xlabel= "Prediction Prob: %f"%(prediction_prob[0,i]))
```

```
[68]: visualize_mislabelled_images(train_x_sample, train_y_sample.
    ↪ T,predicted_labels_train, prediction_prob_train,dataset = "training")
```

Total Mislabelled Images: 0

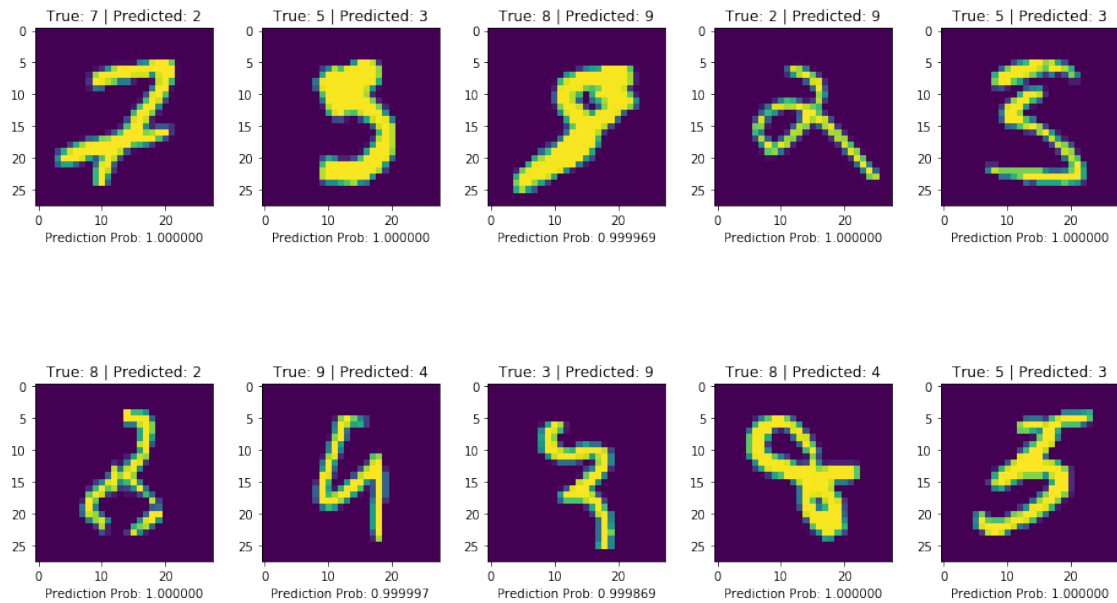
Sample Mislabelled Training Images



```
[69]: visualize_mislabelled_images(dev_x_orig, dev_y_orig.T, predicted_labels_dev,
    ↪ prediction_prob_prob, dataset = "dev")
```

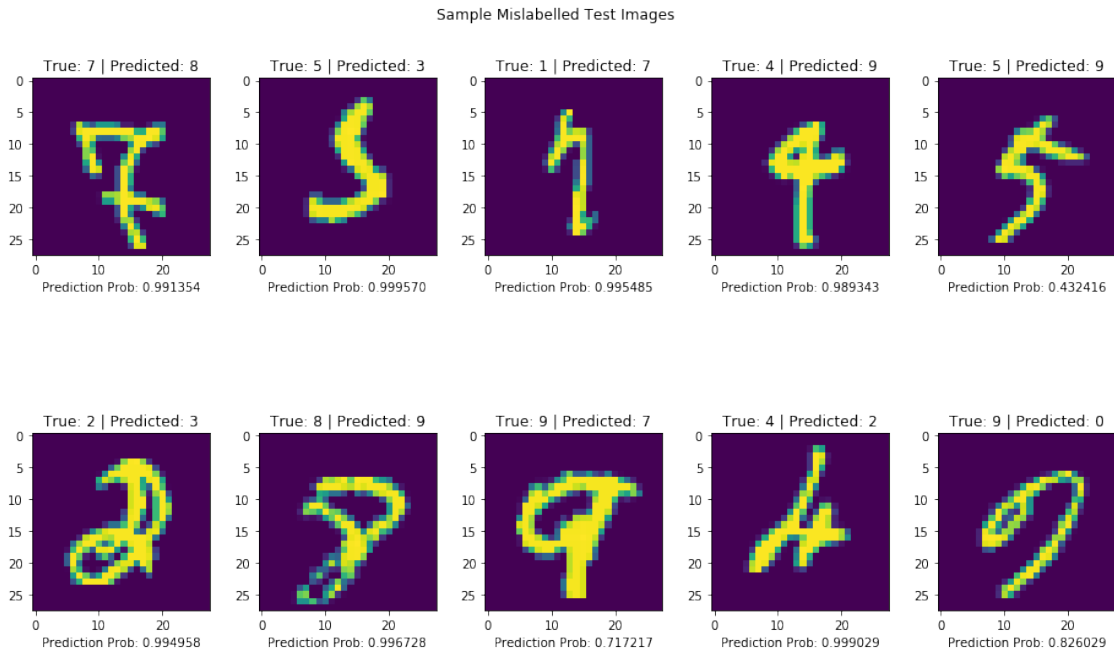
Total Mislabelled Images: 75

Sample Mislabelled Dev Images



```
[70]: visualize_mislabeledled_images(test_x_orig, test_y_orig.T, predicted_labels_test,
    ↪ prediction_prob_prob, dataset = "test")
```

Total Mislabeledled Images: 61



1.4.9 Predicting Real Time images

```
[71]: from PIL import Image
    from dataPrep import one_hot_encoding
```

```
[72]: image_name = "8_1.jpg"
    label = np.array([8]).reshape(1,1)

    fname = "dataset/" + image_name

    image_data = 255 - np.asarray(Image.open(fname).convert('L').resize((28,28)))
    image_flattened = image_data.reshape(image_data.shape[0]*image_data.shape[1],-1)
    image_norm = (image_flattened/255.)


    label_encoded = one_hot_encoding(label)

    predicted_label, pred_prob, acc = predict(image_norm, label_encoded, parameters)

    plt.title("True Label: " + str(label.squeeze()))
```

```
plt.xlabel("Predicted: %d | With Prob: %.4f"%(pridected_label.squeeze(),  
↪pred_prob.squeeze()))  
plt.imshow(image_data, interpolation='nearest', cmap='binary')
```

[72]: <matplotlib.image.AxesImage at 0x7f5e320f19d0>



[]: